

# Exploratory Data Analysis

Importing the [Tanzanian Water Well Data \(https://www.drivendata.org/competitions/7/pump-it-up-data-mining-the-water-table/page/23/\)](https://www.drivendata.org/competitions/7/pump-it-up-data-mining-the-water-table/page/23/)

```
In [1]: import pandas as pd

#Training_set_labels = The dependent variable ['status_group'] for each of the rows in Training_set_values
original_y = pd.read_csv('Training_set_labels.csv')

#Training_set_values = The independent variables for the training set
original_X = pd.read_csv('Training_set_values.csv')

executed in 711ms, finished 18:29:43 2021-03-28
```

## Exploring the Dependent Variables

```
In [2]: original_y.nunique()

executed in 26ms, finished 18:29:45 2021-03-28
```

```
Out[2]: id          59400
status_group      3
dtype: int64
```

```
In [3]: original_y['status_group'].value_counts()

executed in 11ms, finished 18:29:45 2021-03-28
```

```
Out[3]: functional      32259
non functional         22824
functional needs repair   4317
Name: status_group, dtype: int64
```

## Exploring the Independent Variables

```
In [4]: original_X.info()
executed in 71ms, finished 18:29:47 2021-03-28

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 40 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     59400 non-null  int64
1   amount_tsh            59400 non-null  float64
2   date_recorded         59400 non-null  object
3   funder                55765 non-null  object
4   gps_height            59400 non-null  int64
5   installer             55745 non-null  object
6   longitude             59400 non-null  float64
7   latitude              59400 non-null  float64
8   wpt_name              59400 non-null  object
9   num_private           59400 non-null  int64
10  basin                 59400 non-null  object
11  subvillage            59029 non-null  object
12  region                59400 non-null  object
13  region_code           59400 non-null  int64
14  district_code         59400 non-null  int64
15  lga                   59400 non-null  object
16  ward                  59400 non-null  object
17  population            59400 non-null  int64
18  public_meeting        56066 non-null  object
19  recorded_by           59400 non-null  object
20  scheme_management     55523 non-null  object
21  scheme_name           31234 non-null  object
22  permit                56344 non-null  object
23  construction_year     59400 non-null  int64
24  extraction_type       59400 non-null  object
25  extraction_type_group  59400 non-null  object
26  extraction_type_class  59400 non-null  object
27  management            59400 non-null  object
28  management_group      59400 non-null  object
29  payment               59400 non-null  object
30  payment_type          59400 non-null  object
31  water_quality         59400 non-null  object
32  quality_group         59400 non-null  object
33  quantity              59400 non-null  object
34  quantity_group        59400 non-null  object
35  source                59400 non-null  object
36  source_type           59400 non-null  object
37  source_class          59400 non-null  object
38  waterpoint_type       59400 non-null  object
39  waterpoint_type_group  59400 non-null  object
dtypes: float64(3), int64(7), object(30)
memory usage: 18.1+ MB
```

Exploring Smaller DataFrames by Category

```
In [5]: df = original_X.merge(original_y, on='id', how='outer')
len(df.columns)
# Merging independent and dependent dataframes should yield 41 columns.
executed in 190ms, finished 18:29:52 2021-03-28
```

Out[5]: 41

```
In [6]: import matplotlib.pyplot as plt
import numpy as np

def plot_target_bar_graph(category, target, plotname):

    # Function created to plot functionality dependent on any variable in the DataFrame

    # category = df_independent variable
    # target = df_dependent variable (status)
    # plotname = df_independent plot title

    set_labels = category.unique()
    bar_labels = target.unique()
    bar = {}
    for i in range(len(bar_labels)):
        bar[i] = {}
        for j in set_labels:
            bar[i][j] = 0

    count = 0
    for k in category:
        if target[count] == bar_labels[0]:
            bar[0][k] = bar[0][k] + 1
        elif target[count] == bar_labels[1]:
            bar[1][k] = bar[1][k] + 1
        else:
            bar[2][k] = bar[2][k] + 1
        count += 1

    x = np.arange(len(set_labels)) # the label locations
    width = 0.25 # the width of the bars

    fig, ax = plt.subplots()
    ax.bar(x - width/2, list(bar[0].values()), width, label=bar_labels[0], color='royalblue') #bar1
    ax.bar(x, list(bar[1].values()), width, label=bar_labels[1], color='black') #bar2
    ax.bar(x + width/2, list(bar[2].values()), width, label=bar_labels[2], color='lightblue') #bar3

    # Add some text for labels, title and custom x-axis tick labels, etc.
    ax.set_ylabel('Number of Waterpoints')
    ax.set_title(f'Amount of Waterpoints by {plotname}')
    ax.set_xticks(x)
    ax.set_xticklabels(set_labels)
    ax.legend()

    fig.tight_layout()

    plt.show()
```

executed in 232ms, finished 18:30:11 2021-03-28

#### Time Related: date\_recorded, construction\_year

```
In [7]: df_time = df.iloc[:, [2, 23, 40]].copy()
print(df_time.info())
df_time.nunique()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 59400 entries, 0 to 59399
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   date_recorded    59400 non-null  object
1   construction_year 59400 non-null  int64
2   status_group     59400 non-null  object
dtypes: int64(1), object(2)
memory usage: 1.8+ MB
None
```

```
Out[7]: date_recorded    356
construction_year      55
status_group           3
dtype: int64
```


```
In [8]: df_time['date_recorded'] = pd.to_datetime(df_time['date_recorded'], format='%Y-%m-%d')
df_time['year_recorded'] = pd.DatetimeIndex(df_time['date_recorded']).year
df_time.year_recorded.value_counts()
```

executed in 20ms, finished 18:30:14 2021-03-28

```
Out[8]: 2011    28674
2013    24271
2012     6424
2004       30
2002         1
Name: year_recorded, dtype: int64
```

In [9]: `plot_target_bar_graph(df_time.year_recorded, df_time.status_group, 'Year Recorded')`

executed in 991ms, finished 18:30:21 2021-03-28



In [10]: `df_time.construction_year.value_counts()`

executed in 6ms, finished 18:30:23 2021-03-28

Out[10]:

0	20709
2010	2645
2008	2613
2009	2533
2000	2091
2007	1587
2006	1471
2003	1286
2011	1256
2004	1123
2012	1084
2002	1075
1978	1037
1995	1014
2005	1011
1999	979
1998	966
1990	954
1985	945
1980	811
1996	811
1984	779
1982	744
1994	738
1972	708
1974	676
1997	644
1992	640
1993	608
2001	540
1988	521
1983	488
1975	437
1986	434
1976	414
1970	411
1991	324
1989	316
1987	302
1981	238
1977	202
1979	192
1973	184
2013	176
1971	145
1960	102
1967	88
1963	85
1968	77
1969	59
1964	40
1962	30
1961	21
1965	19
1966	17

Name: construction\_year, dtype: int64

```
In [ ]: counter = 0
for year in df_time['construction_year']:
    if year >= 2010:
        df_time.loc[count, 'construction_decade'] = 'After_2010'
    if year >= 2000:
        df_time.loc[count, 'construction_decade'] = '2000-2010'
    if year >= 1990:
        df_time.loc[count, 'construction_decade'] = '1990-2000'
    if year >= 1980:
        df_time.loc[count, 'construction_decade'] = '1980-1990'
    if year >= 1970:
        df_time.loc[count, 'construction_decade'] = '1970-1980'
    if year >= 1960:
        df_time.loc[count, 'construction_decade'] = '1960-1970'
    else:
        df_time.loc[count, 'construction_decade'] = 'Before_1950'
    counter += 1
```

executed in 26.4s, finished 17:48:07 2021-03-28

```
In [ ]: df_time.construction_decade.value_counts()
```

executed in 24.3s, finished 17:48:07 2021-03-28

```
In [ ]: plot_target_bar_graph(df_time.year_recorded, df_time.status_group, 'Year Recorded')
```

executed in 10ms, finished 18:22:42 2021-03-28

The time related information provided will not be further explored in the dataset used for modelling. There is not distinct relationship found in the plots that compare the functionality of the waterpoints versus the time the information was recorded or when the waterpoint was created.

**Geographical Descriptions: gps\_height, longitude, latitude**

```
In [ ]: df_geographical = df.iloc[:, [4, 6, 7, 40]].copy()
df_geographical.head()
```

executed in 20ms, finished 16:46:07 2021-03-28

```
In [ ]: # Intial plot showed an outlier where Long/Lat = 0.
df.loc[df['longitude']==0].head()
```

executed in 92ms, finished 16:57:35 2021-03-28

```
In [ ]: # Dropping outlier for a better plot.
df_geographical = df_geographical[df_geographical.longitude != 0]
```

executed in 11ms, finished 16:57:48 2021-03-28

```
In [ ]: import seaborn as sns

plt.figure(figsize=(10,10))
sns.scatterplot(data=df_geographical,x='longitude', y='latitude', hue='status_group', palette=['green', 'red', 'yellow'])
plt.title('Status of Waterpoints by Location in Tanzania')
plt.show()
```

executed in 6.43s, finished 16:57:57 2021-03-28

```
In [ ]: import seaborn as sns

plt.figure(figsize=(10,10))
sns.scatterplot(data=df_geographical,x='longitude', y='latitude', hue='gps_height')
plt.title('Waterpoints by GPS Height in Tanzania')
plt.show()
```

executed in 6.28s, finished 16:58:03 2021-03-28

The geographical information provided will not be further explored in the dataset used for modelling. The status of the waterpoints seem to be evenly distributed throughout Tanzania. Other location and population data may show more correlation.

**Location Information: subvillage, region, region\_code, district\_code, lga, ward**

```
In [ ]: df_location = df.iloc[:, [11, 12, 13, 14, 15, 16, 40]].copy()
df_location.head()
```

executed in 41ms, finished 16:58:03 2021-03-28

```
In [ ]: df_location.nunique()
```

executed in 70ms, finished 16:58:04 2021-03-28

```
In [ ]:
```

Community/Interaction Information: population, public\_meeting

In [ ]:

df\_community = df.iloc[:, [17, 18, 40]].copy()  
df\_community.head()

executed in 20ms, finished 03:44:37 2021-03-28

In [ ]:

df\_community['public\_meeting'].describe()

executed in 18ms, finished 03:48:14 2021-03-28

Management (Creation -> Current):

In [ ]:

df\_management = df.iloc[:, [3, 5, 20, 21, 27, 28, 40]].copy()  
df\_management.head()

executed in 30ms, finished 20:42:28 2021-03-27

Extraction Type:

In [ ]:

df\_extraction = df.iloc[:, [24, 25, 26, 40]].copy()  
df\_extraction.head()

executed in 21ms, finished 20:46:03 2021-03-27

Source Type:

In [ ]:

df\_source = df.iloc[:, [35, 36, 37, 40]].copy()  
df\_source.head()

executed in 16ms, finished 20:46:55 2021-03-27

Quality of Water:

In [ ]:

df\_quality = df.iloc[:, [31, 32, 40]].copy()  
df\_quality.head()

executed in 20ms, finished 20:47:41 2021-03-27

Quantity of Water:

In [ ]:

df\_quantity = df.iloc[:, [1, 33, 34, 40]].copy()  
df\_quantity.head()

executed in 22ms, finished 20:48:20 2021-03-27

Waterpoint Type:

In [ ]:

df\_waterpoint = df.iloc[:, [38, 39, 40]].copy()  
df\_waterpoint.head()

executed in 19ms, finished 20:49:22 2021-03-27

Misc. Information:

In [ ]:

df\_misc = df.iloc[:, [0, 8, 9, 10, 22, 29, 30, 40]].copy()  
df\_misc.head()

executed in 36ms, finished 20:51:33 2021-03-27

In [ ]:

In [ ]:

In [ ]:

*#Transforming data for analysis purpose*  
  
count = 0  
original\_y['status\_group\_binary'] = ""  
original\_y['status\_group\_tertiary'] = ""  
  
for f in original\_y['status\_group']:  
 if f == 'functional':  
 original\_y.loc[count, 'status\_group\_binary'] = 1  
 original\_y.loc[count, 'status\_group\_tertiary'] = 1  
 elif f == 'non functional':  
 original\_y.loc[count, 'status\_group\_binary'] = 2  
 original\_y.loc[count, 'status\_group\_tertiary'] = 2  
 else:  
 original\_y.loc[count, 'status\_group\_binary'] = 1  
 original\_y.loc[count, 'status\_group\_tertiary'] = 3  
 count += 1

executed in 1m 41.0s, finished 00:07:14 2021-03-26

```
In [ ]: from pandas.api.types import is_string_dtype
        from pandas.api.types import is_numeric_dtype

        for column in original_X:
            if is_numeric_dtype(column):
                column = column.fillna(0)
            else:
                original_X[column] = original_X[column].fillna('unknown')
```

executed in 108ms, finished 00:07:14 2021-03-26

```
In [ ]: import matplotlib.pyplot as plt

        tracker = []

        for x in original_X:
            s2 = original_X[x]
            if len(s2.value_counts()) < 10:
                tracker.append(x)
                prob = s2.value_counts()
                prob.plot(kind='bar')
                plt.xticks(rotation=90)
                plt.title(x)
                plt.show()
            else:
                pass
```

executed in 2.23s, finished 00:07:16 2021-03-26

```
In [ ]: df = original_X[tracker]
```

executed in 13ms, finished 00:07:16 2021-03-26

```
In [ ]: df['functions'] = original_y['status_group_binary']
```

executed in 23ms, finished 00:07:16 2021-03-26

```
In [ ]: import numpy as np
        from sklearn.model_selection import train_test_split
        from sklearn.impute import SimpleImputer
        from sklearn.preprocessing import StandardScaler
        from sklearn.preprocessing import OneHotEncoder

        def preprocess(X, y):
            '''Takes in features and target and implements all preprocessing steps for categorical and continuous features returning
            train and test DataFrames with targets'''

            # Train-test split (75-25), set seed to 10
            X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10)

            # OneHotEncode Categorical variables
            ohe = OneHotEncoder(handle_unknown='ignore')

            X_train_ohe = ohe.fit_transform(X_train)
            X_test_ohe = ohe.transform(X_test)

            columns = ohe.get_feature_names(input_features=X_train.columns)
            cat_train_df = pd.DataFrame(X_train_ohe.todense(), columns=columns)
            cat_test_df = pd.DataFrame(X_test_ohe.todense(), columns=columns)

            # Combine categorical and continuous features into the final dataframe
            X_train_all = pd.concat([cat_train_df], axis=1)
            X_test_all = pd.concat([cat_test_df], axis=1)

            return X_train_all, X_test_all, y_train, y_test
```

executed in 819ms, finished 00:07:17 2021-03-26

```
In [ ]: y = df['functions']
        X = df.drop(['functions'], axis=1)

        X_train_all, X_test_all, y_train, y_test = preprocess(X, y)
```

executed in 189ms, finished 00:07:17 2021-03-26

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

